

Ninad Parthiv Shah

+1-602-884-6345 | ninadpshah2@gmail.com | [linkedin.com/ninadpshah](https://www.linkedin.com/in/ninadpshah) | github.com/NinadPShah | [Portfolio Website](#)

Professional Summary

AI Engineer skilled in Python, LLMs, and prompt engineering with expertise in building AI workflows, Claude Code driven automation, and SQL pipelines that prepare data for AI consumption and business reporting. Strong foundation in agentic systems, RAG pipelines, and BI dashboards, with proven ability to partner cross functionally with stakeholders and translate business needs into data driven solutions.

Skills Summary

- **Languages:** C, C++, Python, SQL
- **AI/ML:** LangChain, LangGraph, Google ADK, Claude Agent SDK, PyTorch, Scikit-Learn, Pandas, NumPy, Hugging Face, mlflow
- **Tools:** BigQuery, AWS, Docker, FastAPI, REST APIs, PySpark, Spark, Hadoop, Azure, Snowflake, Git, Jupyter, Tableau, Power BI
- **Modelling Techniques:** Multi-Agent Systems, Agent Orchestration, MCP Servers, RAG Pipelines, Retrieval Pipelines, Evaluation Pipelines, Causal Inference, Predictive Modeling, Feature Engineering, Time Series Analysis, LLMs, Prompt Engineering, CI/CD
- **Certifications:** AWS Cloud Practitioner ([link](#)), AWS AI Practitioner ([link](#)), Hadoop-Apache Spark ([link](#)), Microsoft Azure ([link](#)),

Education

Arizona State University, Tempe, Arizona

Aug 2023 – May 2025

Master of Science, Data Science, Analytics, and Engineering (GPA: 3.97)

- **Coursework:** Statistical Machine Learning, Data Visualization, Data-Driven Optimization, Data Processing at Scale(Big Data)

Indian Institute of Technology Gandhinagar, India

July 2019 – July 2023

Bachelor of Technology, Electrical Engineering, Computer Science

Professional Experience

Armor1.ai

Dec 2025 – Present

AI Security Software Engineer

Remote

- Designed a 5-layer AI workflow with multi-pass LLM analysis and prompt engineering across 25+ metrics, version controlling prompts and tuning guardrails to optimize accuracy and cut per scan cost to under \$0.005.
- Built SQL transformation pipelines over BigQuery telemetry to clean and structure data for AI consumption, adding confidence scoring and a 10-layer false-positive filter to surface trusted findings.
- Generated reports, KPIs, and dashboards for leadership, turning raw telemetry into prioritized guidance for non-technical stakeholders.

CAPX

Aug 2025 – Dec 2025

AI Software Engineer

Remote

- Built AI agents using Google ADK and LangGraph with sequential and autonomous modes, extracting 20+ metrics from 4000+ firms.
- Built backend services with multi-step workflow automation, prompt engineering, and tool-calling patterns, adding state management and guardrails for usage analysis.
- Owned event-driven architectures to ship AI-powered tools, prototyping with Claude Code and iterating on real world feedback.

Arizona State University

Sep 2024 – May 2025

Academic Integrity Data Analyst

Tempe, AZ

- Defined project scope and deliverables processing 500+ hours of data, building solutions with attention to detail and fast iteration.
- Communicated complex technical findings clearly to stakeholders at different levels through well-documented deliverables.

Spyne.ai

May 2022 – Jul 2022

Computer Vision Intern

Gurugram, India

- Built and deployed production AI services using FastAPI and AWS cloud infrastructure, achieving 99.8% uptime through robust testing.
- Designed computer vision system trained on 50K+ images, rapidly prototyping and iterating based on real customer feedback.
- Integrated AI services into production applications via REST APIs, implementing logging, telemetry, monitoring, and CI/CD pipelines.
- Fine-tuned and deployed image enhancement models into production APIs with 95%+ accuracy, ensuring reliability and performance.

Academic Projects

IPO Outcome Modeling from SEC EDGAR Filings

- Built models on 16 years of SEC EDGAR data for 9,000+ companies, engineering features to link financials to IPO outcomes.
- Used Python, pandas, AWS S3, and PostgreSQL to build feature pipelines and storage for large-scale modeling workflows.
- Applied causal inference and statistical testing to isolate outcome drivers, validating predictive models with cross-validation on data.

SEC EDGAR RAG Pipeline Local MCP Server

- Built a production MCP server with a RAG pipeline over SEC EDGAR 10-K filings using FastMCP and ChromaDB, exposing tool-calling endpoints across 15+ filings with metadata-filtered semantic search.
- Implemented Loughran-McDonald sentiment scoring across 5 financial categories to automate risk factor analysis across filings.
- Developed year-over-year filing comparison tool detecting material language changes, grounded in the Lazy Prices research.